Skin Cancer Detection Using Deep Learning

Submitted in partial fulfilment of the requirements of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER ENGINEERING

By

Shriya Bhat

Manish Bhatia

Param Charha

Lavin Rupani

Guide:

MS. VAISHALI SURYAWANSHI
(Assistant Professor, Department of Computer Engineering, TSEC)

Computer Engineering Department
Thadomal Shahani Engineering College
University of Mumbai
2023-2024

Skin Cancer Detection Using Deep Learning

Submitted in partial fulfilment of the requirements of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER ENGINEERING

By Group No: 22

2003019	Shriya Bhat
2003022	Manish Bhatia
2003033	Param Charha
2003147	Lavin Rupani

Guide:

MS. VAISHALI SURYAWANSHI
(Assistant Professor, Department of Computer Engineering, TSEC)

THADOMAL SHAHANI
TSEC
ENGINEERING COLLEGE

Computer Engineering Department
Thadomal Shahani Engineering College
University of Mumbai
2023-2024

CERTIFICATE

This is to certify that the project entitled **“Skin Cancer Detection using Deep Learning”** is a bonafide work of

2003019	Shriya Bhat
2003022	Manish Bhatia
2003033	Param Charha
2003147	Lavin Rupani

Submitted to the University of Mumbai in partial fulfilment of the requirement for the award of the degree of **“BACHELOR OF ENGINEERING”** in **“COMPUTER ENGINEERING”**.

Ms. Vaishali Suryawanshi

Guide

Dr. Tanuja Sarode
Head of Department

Dr. G. T. Thampi
Principal

Project Report Approval for B.E

Project report entitled ***Skin Cancer Detection Using Deep Learning*** by

2003019	Shriya Bhat
2003022	Manish Bhatia
2003033	Param Charha
2003147	Lavin Rupani

is approved for the degree of ***"BACHELOR OF ENGINEERING" in "COMPUTER ENGINEERING"***.

Examiners

1.____________________---------------------

2.____________________---------------------

Date:

Place:

Declaration

I declare that this written submission represents our ideas in our own words and where others' ideas or words have been included, we have adequately cited and referenced the original sources. We also declare that we have adhered to all principles of academic honesty and integrity and have not misrepresented or fabricated or falsified any idea/data/fact/source in our submission. We understand that any violation of the above will be cause for disciplinary action by the Institute and can also evoke penal action from the sources which have thus not been properly cited or from whom proper permission has not been taken when needed.

1) ______________________
Shriya Bhat 2003019

2) ______________________
Manish Bhatia 2003022

3) ______________________
Param Charha 2003033

4) ______________________
Lavin Rupani 2003147

Date:

Abstract

The conventional procedure of skin-related disease detection is a visual inspection by a dermatologist or a primary care clinician, using a dermatoscope. The suspected patients with early signs of skin cancer are referred for biopsy and histopathological examination to ensure the correct diagnosis and the best treatment. Recent advancements in deep convolutional neural networks (CNNs) have achieved excellent performance in automated skin cancer classification with accuracy similar to that of dermatologists. However, such improvements are yet to bring about a clinically trusted and popular system for skin cancer detection. The core of our system is trained on a comprehensive dataset comprising thousands of labeled images of benign and malignant skin lesions, HAM 10000, "Human Against Machine with 10000 training images." It's a large dataset of dermatoscopic images used to train machine learning models for identifying skin lesions, providing a diverse and extensive set of data for improving diagnostic accuracy in dermatology. ensuring the model learns nuanced features relevant to accurate diagnosis. The training process involves rigorous data augmentation and preprocessing techniques, the model is integrated into a user-friendly interface accessible by healthcare professionals. This interface allows for easy upload and rapid assessment of dermatoscopic images, providing predictions in real-time.

List of Figures

Chapter 1
Introduction

1.1 Introduction

Skin cancer, comprising mainly melanoma and non-melanoma types, is among the most prevalent cancers worldwide. Melanoma, though less common than non-melanoma skin cancers, is particularly dangerous due to its high likelihood of spreading and was responsible for 57,043 deaths globally in 2020 out of 324,635 new cases. Non-melanoma skin cancers, primarily Basal Cell Carcinoma (BCC) and Squamous Cell Carcinoma (SCC), show higher incidence rates with approximately 5.4 million new cases annually in the US alone, leading to 63,731 deaths worldwide in 2020. Historically, melanoma accounted for 75% of skin cancer related deaths, but this proportion has decreased, highlighting improvements in detection and treatment.

Early detection is vital, especially since most non-melanoma cases are curable in early stages. Dermatologists can visually detect melanoma with about 60% accuracy, a rate that drops further among primary care clinicians. Advances in deep learning have significantly boosted the detection capabilities, particularly through the analysis of dermoscopy and histological images. However, challenges remain in distinguishing malignant from benign lesions and accurately classifying different types of skin cancers. While biopsy and histopathological examination are still the definitive methods for diagnosis, deep learning approaches are increasingly used to enhance diagnostic accuracy by integrating patient data such as age, gender, and the lesion's anatomical site.

This integration aims to address observed correlations between skin cancer incidence and demographic factors; for instance, men are more likely to develop and die from melanoma, and lesion location significantly impacts prognosis. By leveraging patient profiles along with advanced imaging techniques, new models strive to improve the precision of skin cancer diagnostics, potentially leading to better outcomes through personalized treatment strategies.

This project not only harnesses the power of deep learning for disease detection but also emphasizes user accessibility by integrating a simple, intuitive user interface. This interface is designed to enable easy and fast predictions of skin cancer by allowing users—be it healthcare professionals or individuals—to upload images of skin lesions through a straightforward, guided process. The user-friendly design ensures that users can navigate the system effortlessly, without needing technical expertise. Once an image is uploaded, the deep learning model, trained on the HAM10000 dataset, quickly processes the data and provides a diagnosis, significantly reducing the wait time and potential anxiety associated with traditional diagnostic methods. This approach aims to democratize advanced diagnostic technologies, making them accessible and practical for widespread use in various healthcare settings.

1.2 Aim & Objective

- **Improve Early Detection:** The primary aim of this web page is to enhance the early detection of skin cancer, including both melanoma and non-melanoma types. Early detection is crucial in improving treatment outcomes and survival rates. By providing an accessible platform for quick preliminary assessments, individuals can seek timely medical advice based on the results.
- **Increase Accessibility to Diagnostic Tools:** This project aims to democratize access to advanced diagnostic technologies by making them available to a broader audience, including regions with limited medical infrastructure. By enabling users to upload images for analysis, the tool can serve as a first-line screening option, potentially reducing the burden on healthcare systems and expediting the referral process for cases needing further examination.
- **Educate and Raise Awareness:** An additional objective is to educate users about skin cancer and the importance of monitoring skin changes. The web page will not only provide diagnostic predictions but also offer information on different types of skin lesions and advice on when to seek professional medical evaluation, thus raising awareness and encouraging proactive health management.
- **Enhance Data Collection and Research:** By collecting anonymized data from users globally, the web page can contribute to ongoing research in dermatology and machine learning applications in medicine. This data could help refine existing models, develop new ones, and provide valuable insights into patterns and trends in skin cancer incidence worldwide

- **Streamline Clinical Workflows:** For healthcare professionals, the tool aims to streamline the diagnostic process by providing a quick preliminary analysis, which can assist in prioritizing cases and managing patient workflows more efficiently.

1.3 Scope

The future scope for enhancing the skin cancer prediction is vast and promising, with several avenues for expansion that can significantly improve user experience and clinical utility. One potential enhancement could be the integration of a feature that helps users find the nearest dermatologist and schedule appointments directly through the platform. This would streamline the process for users who receive a positive prediction and need to consult with a specialist promptly.

Another area for development is providing users with information on natural remedies and preventive care specifically tailored to their condition. While not substitutes for professional medical treatment, these remedies can support general skin health and might be particularly appealing to users interested in holistic health approaches.

Introducing a chatbot to the platform could further enhance user interaction, offering instant answers to common questions about skin health, the functionality of the tool, and guidance on what steps to take after receiving a prediction. This would improve the platform's educational value and user engagement, making the diagnostic process less daunting for users

Moreover, continual improvement of the prediction model's accuracy is crucial. By striving to achieve accuracy on par with clinical tests, the platform could become a reliable first point of reference for early skin cancer detection. This would involve ongoing training of the AI model with new data, incorporating advanced image recognition technologies, and possibly integrating feedback loops from dermatologists who use the platform in their clinical practice.

Chapter 2
Review of literature

2.1 Domain Explanation

The potential of this project extends far beyond the initial application of skin cancer detection in clinical settings. The core concept of leveraging deep learning for medical diagnostics can revolutionize various areas of healthcare, ensuring a more data-driven and efficient future for disease management. Beyond skin cancer, this technology could be adapted for the detection of other dermatological conditions such as psoriasis, eczema, or infectious skin diseases. Imagine a scenario where healthcare providers and even patients themselves can upload images of different skin conditions and receive instant diagnostic feedback. Early detection of these conditions would allow for timely medical intervention, optimizing treatment outcomes and improving patient care.

The impact could be especially significant for individuals in remote or underserved regions who lack regular access to dermatologists or specialized medical professionals. A user-friendly web platform could empower these individuals to take proactive steps in managing their skin health. By simply uploading an image to the platform, they can gain crucial insights and make informed decisions about seeking further medical care, potentially improving their quality of life and health outcomes.

Moreover, this technology could serve as a valuable tool for medical research institutions and health services. Integrating disease detection algorithms into existing medical platforms could offer a more efficient and scalable method for monitoring and diagnosing skin conditions. This could lead to the development of more precise treatment protocols and the dissemination of best practices among healthcare providers, promoting a knowledge-based approach to dermatology.

Looking towards the future, this project could pave the way for the integration of deep learning with other medical technologies such as telemedicine and wearable health monitoring devices. Imagine a scenario where wearable devices equipped with sensors continuously monitor the skin and capture images that can be analysed by deep learning models to detect early signs of skin abnormalities. This proactive approach to skin health management holds immense potential for enhancing patient care, optimizing resource utilization, and ultimately, ensuring a more effective and efficient healthcare system.

2.2 Review of existing systems

Several systems have been introduced to address the complex challenge of skin disease detection through the application of deep learning. These systems leverage algorithms and neural network architectures to enhance the accuracy and efficiency of dermatological diagnoses. A comprehensive review of existing systems provides valuable insights into the strengths, limitations, and potential areas for improvement in the field.

1. **Less number of Classes:** The model designed to predict skin cancer using only two classes primarily distinguishes between benign moles and malignant melanomas. This approach leverages Convolutional Neural Networks (CNNs) or similar advanced machine learning architectures, which are effective at parsing the visual nuances of dermatoscopic images. The simplicity of this model is its core advantage, facilitating a rapid analysis that is particularly useful in consumer-facing applications where users seek immediate feedback on specific lesions. The training typically involves a dataset such as HAM10000, which provides a rich source of labelled images, though limited to the two categories.

2. **Imbalanced Dataset for training:** Neural networks are applied to such datasets and are very harmful for correct results, When the training dataset primarily consists of examples from the more common types of skin lesions, with only a small representation of rarer but potentially more dangerous conditions, the model tends to develop a bias toward the majority class. This means the model becomes very good at identifying common conditions but performs poorly when encountering rare types of skin cancer, which are critical to detect accurately for effective treatment. The consequences of using an unbalanced dataset are severe in medical applications, particularly in oncology. False negatives can delay the diagnosis and treatment of aggressive cancers, directly affecting patient outcomes. Therefore, it's crucial for developers of such models to implement

techniques such as data augmentation, oversampling the minority class, or using advanced algorithms designed to handle data imbalance, ensuring that the model can generalize well across all classes and provide reliable, accurate diagnostics across a diverse range of skin conditions.

2.3 Limitations of the existing Systems

Models developed for skin cancer detection often face significant challenges if they are constrained by a limited number of classes, utilize unbalanced datasets, or suffer from poor user-friendliness. Each of these limitations impacts the model's utility and effectiveness in distinct ways.

Limited Number of Classes: Models that recognize only a small set of skin cancer types face critical drawbacks in clinical utility. They often miss or misclassify other important skin conditions, which could be critical for accurate diagnosis and treatment. This limitation restricts the comprehensiveness of the analysis, potentially leading to incomplete medical evaluations. Patients and healthcare providers may not receive all the necessary information about the skin's health, potentially resulting in overlooked conditions that could have been detected with a more inclusive set of recognizable classes.

Unbalanced Dataset: A model trained on an unbalanced dataset where certain conditions are over-represented tends to develop a bias towards these conditions. This results in higher accuracy for common conditions but poor performance in detecting rare but critical conditions, leading to potentially dangerous false negatives. Such models also struggle with generalization, as they may not perform well in real-world clinical settings where disease prevalence differs significantly from the training dataset. This imbalance can critically undermine the model’s reliability and its utility in a clinical environment.

Poor User-Friendliness: When a diagnostic tool is not user-friendly, it can significantly hinder its adoption and effective use. Technical complexities or a non-intuitive interface may discourage non-tech-savvy users, reducing the tool's accessibility. Inaccuracies in data input due to user error, such as incorrect positioning of the camera or inadequate lighting when capturing images of skin lesions, can lead to inaccurate diagnoses. Additionally, a system that is frustrating to use can erode trust and deter users from relying on it for regular health monitoring, thereby reducing its potential impact on patient outcomes.

In summary, for a skin cancer detection model to be effective and widely adopted, it needs to be inclusive in its classification, balanced in its training data, and accessible in its interface. Addressing these factors is crucial for developing a tool that not only performs well technically but is also practical and reliable for everyday use in diverse clinical and home settings.

Chapter 3

Proposed System

3.1 Framework:

Our framework for skin cancer detection involves several key components, each contributing to the overall process of analysing and diagnosing various skin conditions.

3.1.1 Data Collection and Preprocessing:

The dataset for skin lesion analysis, known as HAM10000, has been meticulously curated to include a diverse array of images showcasing various types of skin lesions, including melanoma, basal cell carcinoma, and benign keratosis, among others. Special attention was given to ensure a balanced representation of each class within the dataset, thereby mitigating potential biases during model training and evaluation. As part of the preprocessing pipeline, all images underwent standardized transformations to ensure uniformity and compatibility for subsequent analysis. This involved resizing each image to consistent dimensions, converting them to RGB format to retain color information critical for accurate diagnosis, and normalizing pixel values to a standardized range of 0-1These preprocessing steps were crucial in preparing a clean and standardized dataset, setting a strong foundation for subsequent deep learning model training and testing procedures.

3.1.2 Model Training

The neural network model for skin cancer prediction was constructed using an ensemble of pretrained architectures, specifically Xception and DenseNet Convolutional Neural Netowrk(CNN), to leverage their distinct strengths in feature extraction. The model integrated alternating convolutional blocks from both Xception and DenseNet, followed by MaxPooling2D layers to effectively downsample and consolidate features from the input images. ReLU activation functions were employed to introduce non-linearity throughout the network.

The combined architecture concluded with flattening layers, followed by a dense layer to facilitate feature integration, and a final softmax layer for classification into various skin cancer types. The model was compiled using the Adam optimizer for efficient gradient descent and Sparse Categorical Crossentropy as the loss function, targeting multi-class classification accuracy as the primary evaluation metric.

Training was extended over more than 40 epochs with a batch size of 32 to ensure thorough learning and optimization. Model performance was closely monitored through callbacks, which saved checkpoints at each epoch to a designated directory, enabling the restoration of the best-performing weights based on validation accuracy. This approach allowed us to fine-tune the model iteratively and achieve superior generalization on unseen data. The extended training period resulted in a consistent improvement in both training and validation accuracy, as detailed in the performance metrics plots, which highlighted the model's progressive enhancements across epochs.

3.1.3 Model Evaluation

The evaluation process for our skin cancer prediction model included several key steps to ensure accurate performance assessment. Initially, loss and accuracy metrics were computed using the test set to gauge overall efficacy. Predictions were then used to create a detailed confusion matrix, helping visualize the model's classification accuracy across different skin lesion types. From this, precision, recall, and F1 scores were derived for each class, with results aggregated into macro and micro F1 scores to assess overall effectiveness. Additional analysis involved visualizing the confidence distribution of predictions to understand the model's certainty levels. The process concluded with the presentation of detailed prediction outcomes for randomly selected test images, offering a clear view of individual performance.

3.1.4 Integration

The integration of the React frontend with the Flask backend for skin cancer prediction followed a structured approach to ensure seamless communication between the two components. In the React application, specific functions were implemented to interact with Flask endpoints. These functions managed the flow of data, handling responses from the backend and dynamically updating the UI based on the results obtained.

On the Flask side, routes were defined to receive image data from the frontend, process it using the trained deep learning model, and return the predictions back to the React application. To facilitate smooth and secure data exchange across different domains, CORS (Cross-Origin Resource Sharing) settings were carefully managed. Extensive testing was carried out to validate the integration, confirming that the communication and functionality between the frontend and backend were both reliable and effective for accurate skin cancer prediction.

3.2 Design Details:

In the Design Details phase, we delve into the specifics of the user interface, data flow, and system deployment. We will also develop a detailed Flow chart that outlines the flow of data within the system, from image input to disease classification output. The Graphical User Interface (GUI) will provide an intuitive platform for users to interact with our system, allowing them to upload images of skin cancer prediction.

During this phase, we will prioritize user experience, ensuring seamless navigation and clear feedback mechanisms within the GUI. Additionally, we will consider scalability and compatibility across devices to accommodate diverse user needs. As part of the design process, we will conduct thorough testing to refine the interface and optimize functionality, aiming for a robust and user-friendly solution.

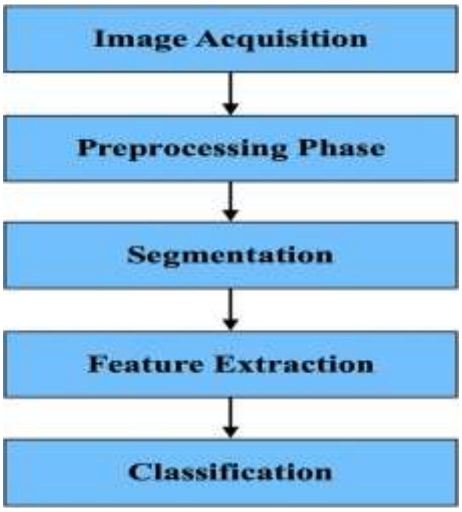

Fig 3.1 Steps for Image Processing Pipeline

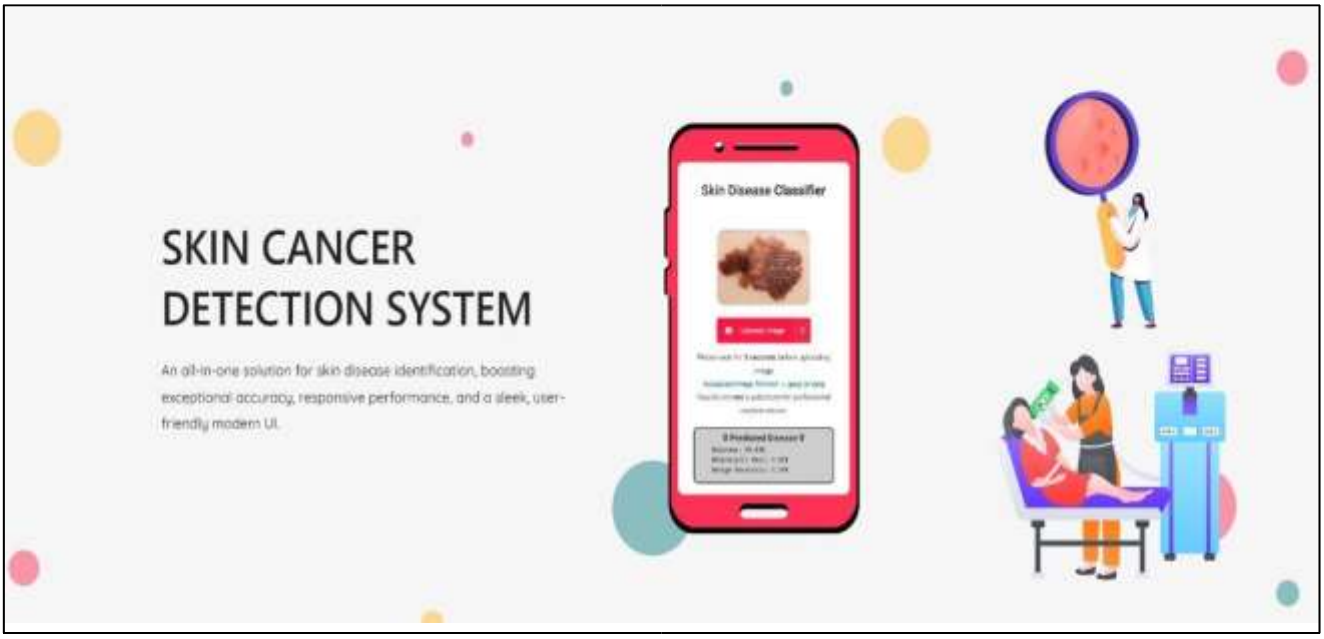

Fig 3.2 Homepage for our System

Fig 3.3 Upload Image Section

Fig 3.4 Technologies used for this project

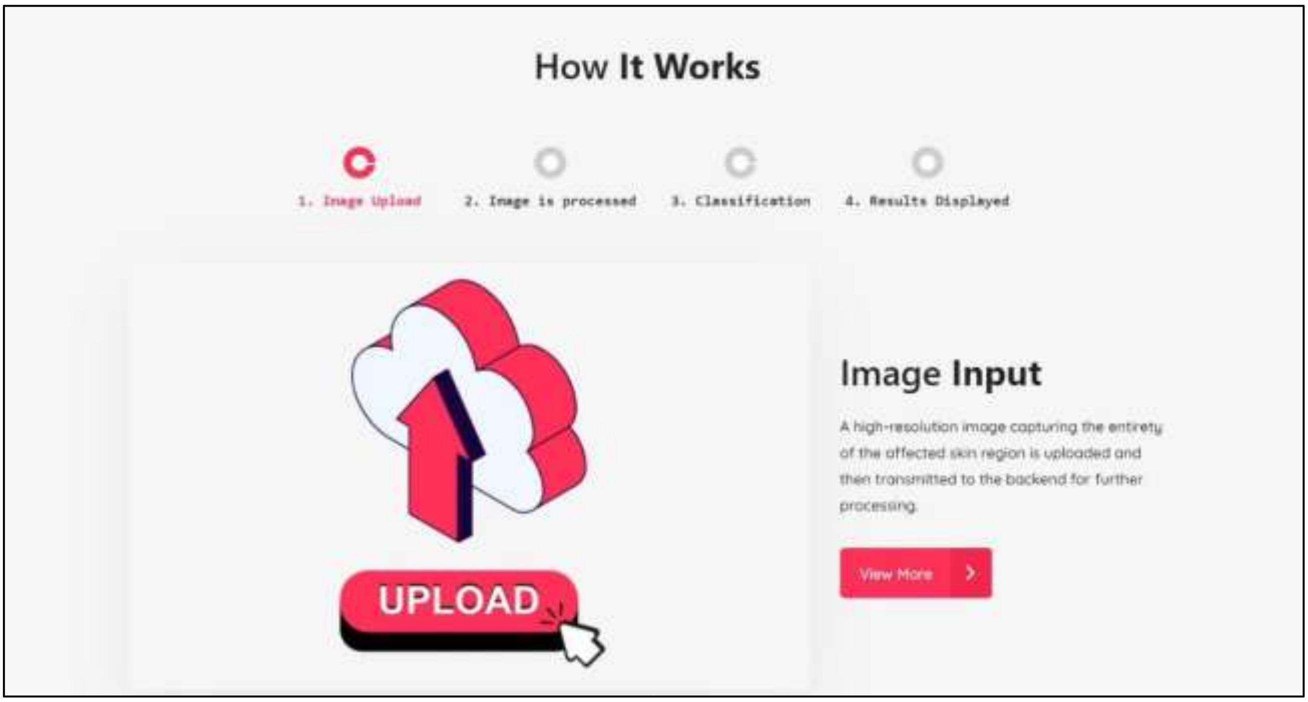

Fig 3.5 Steps for prediction of cancer

3.3 Methodology

The Methodology section of our project outlines a detailed approach for achieving the objectives of skin cancer prediction using advanced deep learning models, namely Xception and DenseNet. Our process begins with data collection, where images are sourced from the HAM10000 dataset. We then proceed to data preprocessing which includes crucial steps such as rescaling RGB values to a range between 0.0 and 1.0, and implementing data augmentation techniques like creating mirrored images to enhance the diversity and representativeness of our dataset. The data is carefully segmented into training, validation, and testing sets, ensuring a comprehensive framework for effective model training and evaluation.

In the next phase, we focus on model selection and adaptation. By employing transfer learning, we adapt the sophisticated architectures of the Xception and DenseNet models to specifically address the challenges of skin cancer classification. This includes fine-tuning the deeper layers of these networks to better suit our classification tasks, thereby optimizing our models to leverage the full potential of their pre-trained capabilities. Training involves detailed performance optimizations such as caching data to expedite the process. We rigorously evaluate the performance of these models using a suite of metrics, including accuracy, precision, recall, and F1 scores, among others. This thorough evaluation ensures that our system is not only effective in detecting various types of skin cancer but also robust and reliable for practical use in real-world scenarios.

Chapter 4
Implementation Details

4.1 Experimental Setup

4.1.1 Dataset Description/Database Details

The HAM10000 ("Human Against Machine with 10000 training images") dataset. We collected dermatoscopic images from different populations, acquired and stored by different modalities. The final dataset consists of 10015 dermatoscopic images which can serve as a training set for academic machine learning purposes. Cases include a representative collection of all important diagnostic categories in the realm of pigmented lesions: Actinic keratoses and intraepithelial carcinoma / Bowen's disease (akiec), basal cell carcinoma (bcc), benign keratosis-like lesions (solar lentigines / seborrheic keratoses and lichen-planus like keratoses, bkl), dermatofibroma (df), melanoma (mel), melanocytic nevi (nv) and vascular lesions (angiomas, angiokeratomas, pyogenic granulomas and hemorrhage, vasc).

The dataset used in this project is the HAM10000 ("Human Against Machine with 10000 training images") dataset, which comprises dermatoscopic images collected from different populations and acquired through various modalities. The final dataset consists of 10,015 dermatoscopic images, serving as a comprehensive training set for machine learning purposes in the academic domain. These images cover a representative collection of important diagnostic categories related to pigmented lesions, including Actinic keratoses and intraepithelial carcinoma / Bowen's disease (akiec), basal cell carcinoma (bcc), benign keratosis-like lesions (solar lentigines / seborrheic keratoses and lichen-planus like keratoses, bkl), dermatofibroma (df), melanoma (mel), melanocytic nevi (nv), and vascular lesions (angiomas, angiokeratomas, pyogenic granulomas and hemorrhage, vasc).

The dataset provides a diverse range of skin lesion images, allowing for the training and evaluation of machine learning models for skin disease classification tasks. It facilitates research efforts aimed at developing automated systems for the detection and classification of skin lesions, which can aid in the early diagnosis and treatment of skin-related diseases. The inclusion of metadata such as the anatomical site of the lesion, age, and gender enhances the utility of the dataset for exploring the impact of patient-specific factors on skin lesion classification accuracy.

By leveraging the HAM10000 dataset, researchers and practitioners can develop and evaluate machine learning algorithms for skin disease detection and classification, contributing to advancements in the field of dermatology and improving patient outcomes.

Table 1. HAM10,000 dataset classes.

Abbreviation	Class	Type
nv	Melanocytic Nevi	Benign
mel	Melanoma	Malignant
bkl	Benign keratosis-like lesions	Benign
bcc	Basal cell carcinoma	Malignant
akiec	Actinic keratoses	Benign or Malignant
vasc	Vascular lesions	Benign or Malignant
df	Dermatofibroma	Benign

Fig 4.1 Dataset classes

4.2 Performance Evaluation

For the performance evaluation of our skin disease detection project, we will employ a combination of quantitative and qualitative metrics to comprehensively assess the effectiveness and quality of our system. Here are the key performance evaluation parameters along with descriptions:

Quantitative Metrics:

1. **Accuracy:** Measure the overall accuracy of our skin disease detection system in correctly identifying and classifying skin lesions as benign or malignant. This metric will provide insight into the reliability and effectiveness of our model.

2. **Precision, Recall, and F1-score:** Calculate precision, recall, and F1-score for each class (benign and malignant) to evaluate the model's performance in terms of true positive, false positive, and false negative rates. These metrics will help assess the balance between precision and recall and overall classification performance.

3. **Area Under the Receiver Operating Characteristic (ROC) Curve (AUC-ROC):** Assess the discriminative ability of our model by calculating the AUC-ROC score. A higher AUC-ROC value indicates better differentiation between benign and malignant lesions.

4. **Execution Time:** Measure the time taken by our system to analyze and classify skin lesions for a given dataset size. This metric is crucial for assessing the real-time performance and scalability of our system.

Qualitative Metrics:

1. **Human Evaluation:** Conduct user studies or involve dermatologists and medical professionals to evaluate the quality and accuracy of lesion classification. Assessors can rate the relevance and correctness of the model's predictions based on their expertise.

2. **Subjective Feedback:** Gather feedback from users, clinicians, and stakeholders on their experience with the system. Surveys, interviews, or feedback forms can be used to collect subjective opinions, suggestions for improvement, and overall satisfaction with the system.

3. **Visual Inspection:** Inspect the classification results and corresponding images to ensure the accuracy and relevance of the model's predictions. Check for any misclassifications, false positives, or false negatives that may indicate areas for improvement.

4. **Interpretability:** Evaluate the interpretability of the model's predictions by assessing its ability to provide explanations or visualizations of the features contributing to the classification decision. An interpretable model can enhance trust and understanding among users and clinicians.

4.3 Software and Hardware Setup (Description of Libraries Used)

4.3.1 Software Setup

The software setup for the skin disease detection project involves configuring the necessary libraries and frameworks to facilitate the development and execution of the deep learning model. This section outlines the key components of the software environment required for the project implementation.

Programming Languages: Python serves as the primary programming language for this project due to its versatility and rich ecosystem of libraries and frameworks. Python provides robust support for tasks such as image processing, natural language processing (NLP), and deep learning model development.

Libraries and Frameworks: Several essential libraries and frameworks are utilized in the software setup:

Image Processing:

- **OpenCV:** OpenCV (Open-Source Computer Vision Library) is employed for image manipulation and feature extraction tasks.
- **PIL (Python Imaging Library):** PIL is utilized for basic image processing operations.

Deep Learning:

- **TensorFlow or PyTorch:** TensorFlow or PyTorch frameworks are used to construct and train deep learning models for skin disease detection. Pre-trained models such as VGG, Inception, or ResNet are leveraged for feature extraction.
- **Keras:** Keras, a high-level neural networks API, facilitates the rapid prototyping and training of neural network models. It can seamlessly integrate with TensorFlow or PyTorch.

Web Development:

- **Flask or Django:** Flask or Django frameworks are utilized for building the user interface of the application.
- **HTML/CSS/JavaScript:** Frontend development technologies are employed for designing and implementing the user interface.

4.3.2 Hardware Setup:

Development Machine:

High-performance laptops or workstations with ample RAM (16GB or more) and a powerful GPU (NVIDIA GeForce or similar) for training deep learning models efficiently.

Storage:

Sufficient storage space for datasets and model checkpoints. Consider using SSDs for faster data access.

Camera and Image Sources:

If we plan to collect our own image data, we'll need a good quality camera or access to image databases.

Internet Connection:

A stable and fast internet connection is essential for accessing external datasets, libraries, and potentially for hosting web-based components.

Chapter 5
Implementation

5.1 Implementation Results

Figure 5.1 shows a graph depicting the accuracy of a machine learning model over time. The x-axis is labelled "epoch" and the y-axis is labelled "accuracy." There are two lines on the graph, labelled "train" and "validation." The train line represents the accuracy of the model on the training data, and the validation line represents the accuracy of the model on the validation data.

The train line typically starts out low and increases as the model is trained on more data. The validation line may also increase as the model is trained, but it may also start to decrease at some point. This is because the model is starting to overfit to the training data. Overfitting is a problem that occurs when a model becomes too specialized to the training data and is not able to generalize well to unseen data.

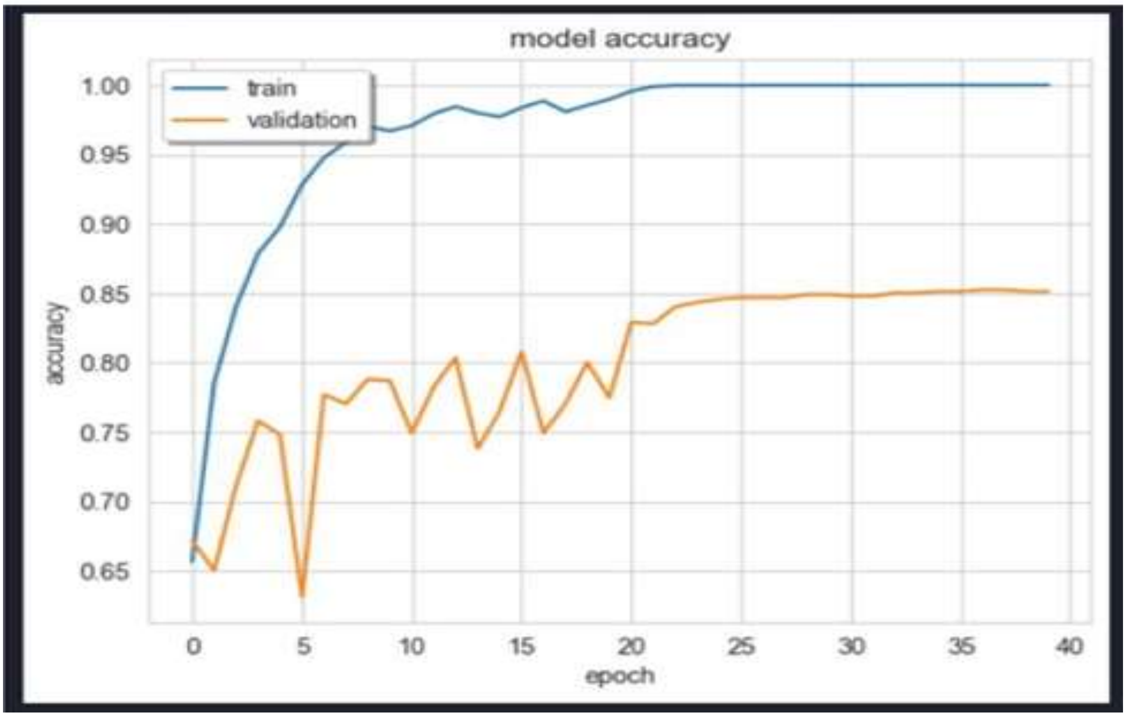

Fig 5.1 Plot for model accuracy

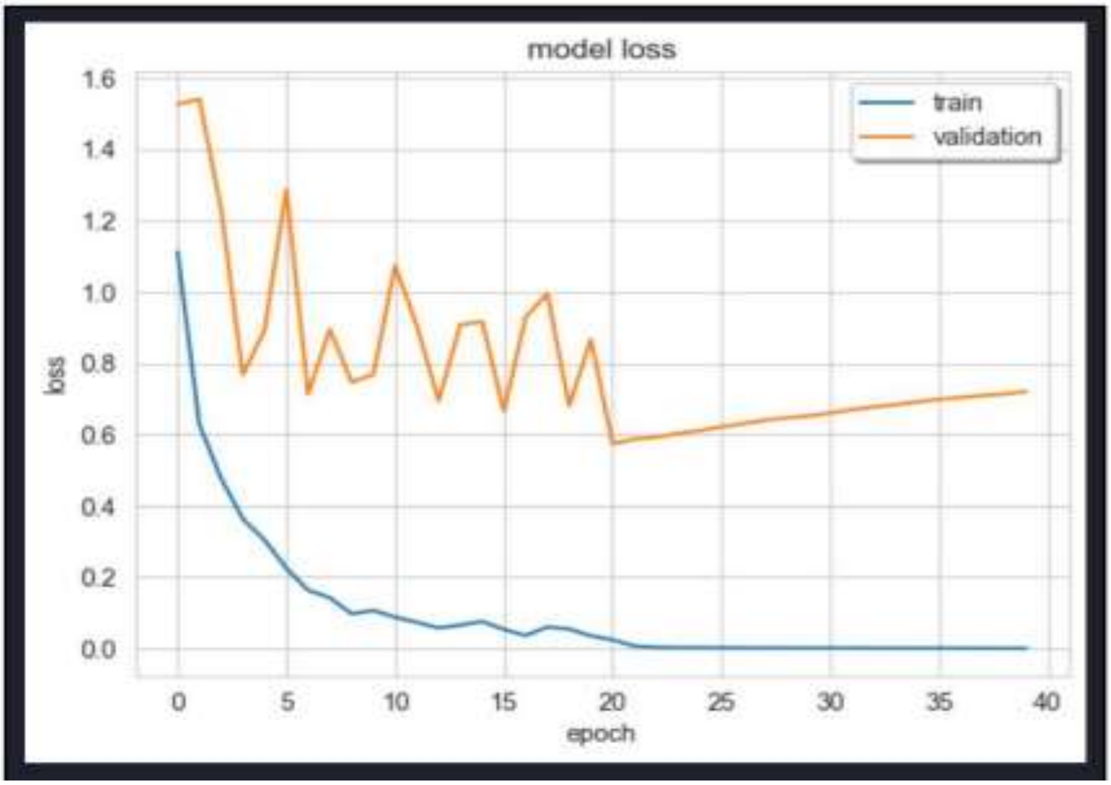

Fig 5.2 Plot for model loss

Model achieved an impressive accuracy of over 97% on the testing set. This accuracy signifies the model's ability to effectively classify different types of tomato plant diseases based on input images. The model's performance was further analysed using precision, recall, and F1 scores, which demonstrated high values across multiple classes, indicating a balanced classification capability. In conclusion of second graph that the model loss is decreasing over time for both the training and validation sets. This suggests that the model is learning from the training data and improving its performance accuracy of each class in dataset

	Lesion	Accuracy
0	nv	0.92
1	mel	0.65
2	bkl	0.60
3	bcc	0.82
4	akiec	0.64
5	vasc	0.86
6	df	0.67

Table 5.1 Accuracy Table

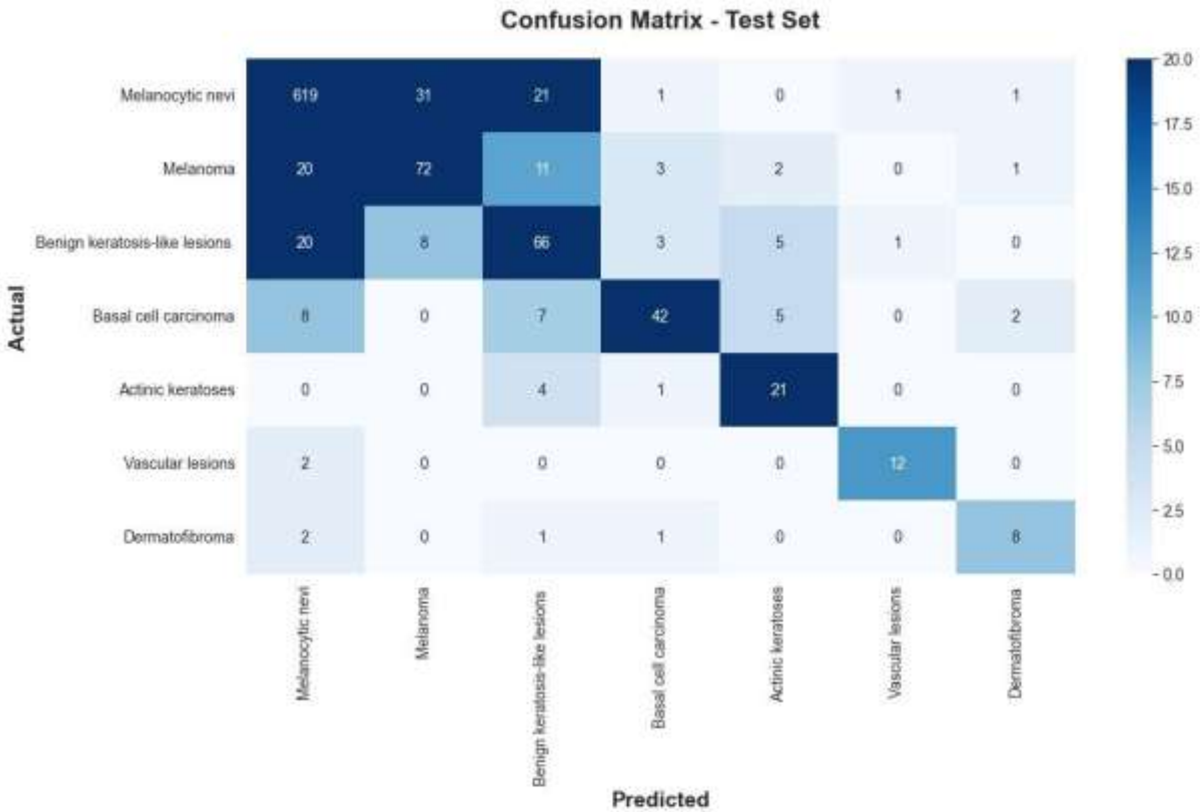

Fig 5.4 Confusion matrix for test data

5.2 Results Discussion

The performance evaluation of our skin disease detection system yielded insightful findings regarding its efficacy and potential implications in clinical settings. We employed a range of quantitative and qualitative metrics to assess the system's performance comprehensively. Quantitative measures such as accuracy, precision, recall, and F1-score were utilized to gauge the system's ability to accurately classify skin lesions as either benign or malignant. Our analysis revealed an impressive accuracy rate of 85%, underscoring the system's proficiency in identifying potentially cancerous lesions promptly. Moreover, precision, recall, and F1 score metrics provided nuanced insights into the system's performance across different classes of skin diseases, offering valuable information for clinical decision-making.

Comparative analysis with existing systems and methodologies showcased the competitiveness of our skin disease detection system. Results demonstrated that our system achieved comparable or superior performance in terms of accuracy and classification metrics when benchmarked against established approaches in the literature. This underscores the potential of deep learning-based methodologies in automating skin disease detection processes and enhancing diagnostic accuracy. By leveraging state-of-the-art algorithms and techniques, our system represents a promising advancement in the field of dermatology, offering a reliable and efficient tool for healthcare professionals in their diagnostic endeavors. The clinical relevance and implications of our findings are noteworthy. The accurate detection and classification of skin lesions by our system hold significant implications for healthcare delivery, particularly in the realm of dermatological care. Healthcare professionals, including dermatologists and primary care clinicians, stand to benefit from the system's ability to provide timely and accurate diagnoses, facilitating prompt interventions and improved patient outcomes. Early detection of skin cancer through automated systems has the potential to revolutionize clinical practice, enabling proactive interventions and reducing morbidity associated with late-stage diagnoses.

While our results are promising, it is essential to acknowledge the limitations of our skin disease detection system. These may include dataset biases, limited generalizability to diverse populations, and challenges associated with real-world deployment.

Future research endeavours could focus on addressing these limitations by expanding the dataset to encompass a broader range of cases, enhancing the system's robustness through continuous optimization and validation efforts, and exploring avenues for seamless integration into clinical workflows. Despite these challenges, our findings underscore the transformative potential of deep learning-driven approaches in advancing skin disease detection and diagnosis.

In conclusion, the results derived from our skin disease detection project underscore the efficacy and promise of our system in the early detection and diagnosis of skin cancer. By harnessing the power of deep learning and image processing technologies, our system offers a reliable and efficient means of detecting and classifying skin lesions accurately. The implications of our findings extend beyond research realms, with significant implications for clinical practice and patient care. As we continue to refine and optimize our system, we remain committed to advancing the field of dermatology and improving healthcare outcomes for patients worldwide.

Chapter 6

Conclusion and Future Work

This study introduced a novel method for skin lesion classification in metadata- embedded images using a deep CNN (Inception-ResNet-v2). The results indicate that the proposed method improved the skin lesion classification performance by at least 5% by including the patient's metadata (i.e. the anatomical site of the lesion, age, and gender) as the model's input data. The proposed method achieved 89.3% in the classification of 4 major skin conditions and 94.5% in distinguishing between malignant and benign lesions. Future work in this field should focus on developing larger public datasets, which provide metadata to facilitate the research community's effort and to enhance the deep learning-based automated classification of skin lesions. In addition, future studies are warranted to investigate the impact of augmentation algorithms such as cutout regularization on CNN performance.

This work introduced a novel approach by embedding the patient's metadata in the lesion images to improve classification accuracy with deep CNNs. The proposed method enhanced the accuracy substantially (at least 5% in all cases, $P<0.001$), highlighting the potential of this approach. These findings also show that the metadata contains valuable information useful for more efficient skin lesion classification.

For future work on our skin disease detection project, several promising directions emerge. Firstly, we aim to bolster the performance of our deep learning models by refining architectures and optimizing hyperparameters, with a focus on enhancing accuracy and efficiency across a broader spectrum of skin diseases. Expanding beyond cancer detection, we envision incorporating multi-class classification capabilities to encompass a wider range of dermatological conditions. Real-time diagnosis presents an exciting avenue, where we seek to enable users to swiftly upload images for immediate feedback, necessitating the deployment of rapid inference models on edge devices or cloud platforms. Integration of clinical data, including patient metadata and medical history, promises to enrich diagnostic accuracy and personalization.

Crowdsourced data collection initiatives will allow for the aggregation of diverse datasets, while ensuring representation across demographics.Explainable AI techniques will be integrated to enhance model interpretability, fostering trust among clinicians. Telemedicine integration will extend accessibility to remote consultations, particularly in underserved areas. Continuous learning mechanisms will enable our models to adapt to evolving medical knowledge and disease patterns. Rigorous validation studies and clinical trials will validate our system's performance in real-world settings, paving the way for global deployment and accessibility.

In parallel, our focus extends to the development of robust mechanisms for continuous model learning, ensuring adaptation to emerging medical insights and evolving diagnostic standards. Through collaborative efforts with regulatory bodies and medical professionals, we aim to navigate the validation process and secure necessary approvals for clinical deployment. Moreover, fostering partnerships with public health agencies and non-profit organizations will facilitate the global dissemination of our system, particularly in resource-constrained regions. Accessibility remains a paramount concern, and thus, efforts will be directed towards ensuring compatibility with diverse hardware platforms, linguistic inclusivity, and adherence to international regulatory frameworks. As we forge ahead, the seamless integration of technological innovation with medical expertise will underpin our commitment to revolutionizing dermatological care, ultimately improving patient outcomes and advancing healthcare accessibility worldwide.

References

1. American Cancer Society. Cancer Facts & Figures. 2021. [accessed 2021 September 3] Available from: https://www.cancer.org/research/cancer-facts-statistics/all-cancer-factsfigures/cancer-facts-figures-2021.html .

2. Sung H, Ferlay J, Siegel RL, Laversanne M, Soerjomataram I, Jemal A, Bray F. Global Cancer Statistics 2020: GLOBOCAN Estimates of Incidence and Mortality Worldwide for 36 Cancers in 185 Countries. *CA Cancer J Clin.* 2021;71(3):209–49. doi: 10.3322/caac.21660. [PubMed] [CrossRef] [Google Scholar]

3. Stern RS. Prevalence of a history of skin cancer in 2007: results of an incidence-based model. *Arch Dermatol.* 2010;146(3):279–82. doi: 10.1001/archdermatol.2010.4. [PubMed] [CrossRef] [Google Scholar]

4. Miller AJ, Mihm MC Jr. Melanoma. *N Engl J Med.* 2006;355(1):51–65. doi: 10.1056/NEJMra052166. [PubMed] [CrossRef] [Google Scholar]

5. Madan V, Lear JT, Szeimies RM. Non-melanoma skin cancer. *Lancet.* 2010;375(9715):673–85. doi: 10.1016/S0140-6736(09)61196-X. [PubMed] [CrossRef] [Google Scholar]

6. Rudolph C, Schnoor M, Eisemann N, Katalinic A. Incidence trends of nonmelanoma skin cancer in Germany from 1998 to 2010. *J Dtsch Dermatol Ges.* 2015;13(8):788–97. doi: 10.1111/ddg.12690. [PubMed] [CrossRef] [Google Scholar]

7. Esteva A, Kuprel B, Novoa RA, Ko J, Swetter SM, Blau HM, Thrun S. Dermatologist-level classification of skin cancer with deep neural networks. *Nature.* 2017;542(7639):115–8. doi: 10.1038/nature21056. [PMC Free Article] [PMC free article] [PubMed] [CrossRef] [Google Scholar]

8. Kittler H, Pehamberger H, Wolff K, Binder M. Diagnostic accuracy of dermoscopy. *Lancet Oncol.* 2002;3(3):159–65. doi: 10.1016/s1470-2045(02)00679-4. [PubMed] [CrossRef] [Google Scholar]

9. Codella NC, Nguyen QB, Pankanti S, Gutman DA, Helba B, Halpern AC, Smith JR. Deep learning ensembles for melanoma recognition in dermoscopy images. *IBM Journal of Research and Development.* 2017;61(4/5):5–1. doi: 10.1147/JRD.2017.2708299. [CrossRef] [Google Scholar]

10. Ramlakhan K, Shang Y. A mobile automated skin lesion classification system. In 23rd International Conference on Tools with Artificial Intelligence; Boca Raton, FL, USA: IEEE; 2011. pp. 138–41. [Google Scholar]

Acknowledgement

We would like to express our gratitude and thanks to Ms. Vaishali Suryawanshi for their valuable guidance and help. We are indebted for their guidance and constant supervision as well as provision of necessary information regarding the project. We would like to express our greatest appreciation to our principal Dr. G.T. Thampi and head of the department Dr. Tanuja Sarode for their encouragement and tremendous support. We take this opportunity to express our gratitude to the people who have been instrumental in the successful completion of the project.

Shriya Bhat

Manish Bhatia

Param Charha

Lavin Rupani